



**NAMIBIA UNIVERSITY
OF SCIENCE AND TECHNOLOGY**

**FACULTY OF COMPUTING AND INFORMATICS
DEPARTMENT OF INFORMATICS**

MASTERS OF DATA SCIENCE

**Trends in Artificial Intelligence and Machine Learning
(TAI911S)**

GROUP ASSIGNMENT 1

STUDENT NAME: ABINER HAIKIYATA (224002090)

STUDENT NAME: LAHYA SOFIKA (225111004)

LECTURER: Prof Quenum

Import libraries for data processing, modeling, and visualization

using CSV, DataFrames, MLJ, Plots, MLJLinearModels, ScientificTypes, MLJBase, Tables, CategoricalArrays, StatisticalMeasures

1. Explicitly import selected functions to avoid naming conflicts

begin

import DataFrames: select, Not, rename!

import MLJ: machine, fit!, predict, predict_mode

end

Explanation:

- We load necessary packages for reading CSVs, managing dataframes, modeling (MLJ), plotting, and evaluation.
- Explicit imports prevent naming conflicts, especially for select, which exists in multiple packages.

2. Load and clean the food wastage dataset

2. Loading and Cleaning dataset

```
md"""  
### 2. Loading and Cleaning dataset  
"""
```

	Country	Year	Food Category	Total Waste (Tons)	Economic_Loss	Avg Waste per Capita (Kg)	Popula (Milli)
	String15	Int64	String31	Float64	Float64	Float64	Float
1	"Australia"	2019	"Fruits & Vegetables"	19268.6	18686.7	72.69	87.59
2	"Indonesia"	2019	"Prepared Food"	3916.97	4394.48	192.52	1153.9
3	"Germany"	2022	"Dairy Products"	9700.16	8909.16	166.94	1006.1
4	"France"	2023	"Fruits & Vegetables"	46299.7	40551.2	120.19	953.05
5	"France"	2023	"Beverages"	33096.6	36980.8	104.74	1105.4
6	"India"	2024	"Fruits & Vegetables"	11962.9	11196.0	136.21	1311.9
7	"Germany"	2024	"Prepared Food"	45038.7	39191.2	179.27	1349.4
8	"China"	2019	"Fruits & Vegetables"	12791.2	12233.3	90.8	1229.2
9	"UK"	2019	"Meat & Seafood"	14795.6	14347.0	128.91	450.33
10	"India"	2019	"Grains & Cereals"	12118.3	13631.2	141.75	359.26
: more							
5000	"France"	2024	"Bakery Items"	8860.27	7360.38	51.5	879.67

begin

```
food_waste_df = CSV.read("global_food_wastage_dataset.csv", DataFrame)
```

```
# Rename specific columns to remove special characters and spaces
```

```
rename!(food_waste_df,  
         "Economic Loss (Million \$)" => "Economic_Loss",  
         "Household Waste (%)" => "Household_Waste")
```

End

Explanation:

- Reads the dataset and stores it in food waste df.
- Renames columns for cleaner access (e.g., to avoid spaces and special characters).

3. Load the fetal health dataset

```
fetal_health_df =
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	se
1	120	0.0	0.0	0.0	0.0	0.
2	132	0.006	0.0	0.006	0.003	0.
3	133	0.003	0.0	0.008	0.003	0.
4	134	0.003	0.0	0.008	0.003	0.
5	132	0.007	0.0	0.008	0.0	0.
6	134	0.001	0.0	0.01	0.009	0.
7	134	0.001	0.0	0.013	0.008	0.
8	122	0.0	0.0	0.0	0.0	0.
9	122	0.0	0.0	0.002	0.0	0.
10	122	0.0	0.0	0.003	0.0	0.
⋮ more						
2126	142	0.002	0.002	0.008	0.0	0.

```
fetal_health_df = CSV.read("fetal_health.csv", DataFrame)
```

Explanation:

- Reads a dataset that includes features and a categorical label (fetal_health) for classification tasks.

4. Prepare features for the food waste dataset (drop target columns)

```
# -----  
# PROBLEM 1: REGRESSION MODEL  
# -----
```

X_fw =

	Country	Year	Food Category	Total Waste (Tons)	Avg Waste per Capita (Kg)	Population (Million)
1	"Australia"	2019	"Fruits & Vegetables"	19268.6	72.69	87.59
2	"Indonesia"	2019	"Prepared Food"	3916.97	192.52	1153.99
3	"Germany"	2022	"Dairy Products"	9700.16	166.94	1006.11
4	"France"	2023	"Fruits & Vegetables"	46299.7	120.19	953.05
5	"France"	2023	"Beverages"	33096.6	104.74	1105.47
6	"India"	2024	"Fruits & Vegetables"	11962.9	136.21	1311.91
7	"Germany"	2024	"Prepared Food"	45038.7	179.27	1349.45
8	"China"	2019	"Fruits & Vegetables"	12791.2	90.8	1229.29
9	"UK"	2019	"Meat & Seafood"	14795.6	128.91	450.33
10	"India"	2019	"Grains & Cereals"	12118.3	141.75	359.26
: more						
5000	"France"	2024	"Bakery Items"	8860.27	51.5	879.67

```
X_fw = select(food_waste_df, Not([:Economic_Loss, :Household_Waste]))
```

5. Convert string features to numeric codes

```
for col in names(X_fw)  
  if eltype(X_fw[!, col]) <: AbstractString  
    X_fw[!, col] = levelcode.(categorical(X_fw[!, col]))  
  end  
end
```

Explanation:

- X_fw will hold the independent variables for modeling.
- If any column contains strings, we convert them to categorical then to numerical codes (levelcode) for ML modeling.

6. Convert cleaned DataFrame to Matrix and then to MLJ-compatible table

```
X_fw_matrix = 5000x6 Matrix{Float64}:
 2.0  2019.0  5.0  19268.6   72.69   87.59
 9.0  2019.0  8.0   3916.97  192.52  1153.99
 7.0  2022.0  3.0   9700.16  166.94  1006.11
 6.0  2023.0  5.0  46299.7   120.19   953.05
 6.0  2023.0  2.0  33096.6   104.74  1105.47
 8.0  2024.0  5.0  11962.9   136.21  1311.91
 7.0  2024.0  8.0  45038.7   179.27  1349.45
 ⋮
 9.0  2023.0  8.0  39980.4    32.34   27.08
 6.0  2021.0  2.0  47524.7    77.41  1087.46
 2.0  2021.0  2.0  32337.7   194.35  1336.32
 5.0  2018.0  7.0  20641.0    21.04   16.13
 2.0  2021.0  2.0  26566.6   197.14  1086.17
 6.0  2024.0  1.0   8860.27    51.5   879.67
```

Convert to matrix and handle missing values

```
X_fw_matrix = Matrix{Float64}(X_fw)
```

```
X_fw_matrix = Matrix{Float64}(X_fw)
```

```
X_fw_table = MLJBase.table(X_fw_matrix, names=propertynames(X_fw))
```

Explanation:

- Convert the DataFrame to a matrix of Float64 (numeric values).
- Convert that matrix into a MLJ table which is needed for training models.

7. Prepare target variables for regression

X_fw_table =

	Country	Year	Food Category	Total Waste (Tons)	Avg Waste per Capita (Kg)	Population (Million)
1	2.0	2019.0	5.0	19268.6	72.69	87.59
2	9.0	2019.0	8.0	3916.97	192.52	1153.99
3	7.0	2022.0	3.0	9700.16	166.94	1006.11
4	6.0	2023.0	5.0	46299.7	120.19	953.05
5	6.0	2023.0	2.0	33096.6	104.74	1105.47
6	8.0	2024.0	5.0	11962.9	136.21	1311.91
7	7.0	2024.0	8.0	45038.7	179.27	1349.45
8	5.0	2019.0	5.0	12791.2	90.8	1229.29
9	19.0	2019.0	7.0	14795.6	128.91	450.33
10	8.0	2019.0	6.0	12118.3	141.75	359.26
: more						

X_fw_table = MLJBase.table(X_fw_matrix, names=propertynames(X_fw))

y_econ =

Add cell (Ctrl + Enter)

48, 8909.16, 40551.2, 36980.8, 11196.0, 39191.2, 12233.3, 14347.0, 13631.2, .

Prepare targets

y_econ = Float64.(skipmissing(food_waste_df.Economic_Loss)) |> collect

y_household =

[53.64, 30.61, 48.08, 31.91, 36.06, 37.09, 68.93, 59.9, 59.01, 42.42, 31.16, 33.79, 45.43, 3

y_household = Float64.(skipmissing(food_waste_df.Household_Waste)) |> collect

```
y_econ = Float64.(skipmissing(food_waste_df.Economic_Loss)) |> collect
```

```
y_household = Float64.(skipmissing(food_waste_df.Household_Waste)) |> collect
```

Explanation:

- Extracts target values (y_econ, y_household) and removes missing values.

8. Create and train models for Economic Loss and Household Waste

```
model_fw = LinearRegressor(
    fit_intercept = true,
    solver = nothing)

# Create and fit models
model_fw = LinearRegressor()

mach_econ =
  trained Machine; caches model-specific representations of data
  model: LinearRegressor(fit_intercept = true, ...)
  args:
    1: Source @255 ↗ ScientificTypesBase.Table{AbstractVector{ScientificTypesBase.Continu
    2: Source @639 ↗ AbstractVector{ScientificTypesBase.Continuous}

mach_econ = machine(model_fw, X_fw_table, y_econ) |> fit!

Training machine(LinearRegressor(fit_intercept = true, ...), ...).
Solver: MLJLinearModels.Analytical
iterative: Bool false
max_inner: Int64 200

mach_household =
  trained Machine; caches model-specific representations of data
  model: LinearRegressor(fit_intercept = true, ...)
  args:
    1: Source @703 ↗ ScientificTypesBase.Table{AbstractVector{ScientificTypesBase.Continu
    2: Source @253 ↗ AbstractVector{ScientificTypesBase.Continuous}

mach_household = machine(model_fw, X_fw_table, y_household) |> fit!

Training machine(LinearRegressor(fit_intercept = true, ...), ...).
Solver: MLJLinearModels.Analytical
iterative: Bool false
max_inner: Int64 200
```

```
mach_econ = machine(model_fw, X_fw_table, y_econ) |> fit!
```

```
mach_household = machine(model_fw, X_fw_table, y_household) |> fit!
```

Explanation:

- Create MLJ machines for each regression task and train them using fit!.

9. Prepare fetal health classification dataset

```
# -----  
# PROBLEM 2: CLASSIFICATION MODEL  
# -----
```

X_fetal =

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	se
1	120	0.0	0.0	0.0	0.0	0.
2	132	0.006	0.0	0.006	0.003	0.
3	133	0.003	0.0	0.008	0.003	0.
4	134	0.003	0.0	0.008	0.003	0.
5	132	0.007	0.0	0.008	0.0	0.
6	134	0.001	0.0	0.01	0.009	0.
7	134	0.001	0.0	0.013	0.008	0.
8	122	0.0	0.0	0.0	0.0	0.
9	122	0.0	0.0	0.002	0.0	0.
10	122	0.0	0.0	0.003	0.0	0.
: more						
2126	142	0.002	0.002	0.008	0.0	0.

```
X_fetal = select(fetal_health_df, Not(:fetal_health))  
  
y_fetal =  
  CategoricalArrays.CategoricalVector{Int64, UInt32, Int64, CategoricalArrays.CategoricalValue{I  
  
y_fetal = categorical(fetal_health_df.fetal_health)
```

```
X_fetal = select(fetal_health_df, Not(:fetal_health))
```

```
y_fetal = categorical(fetal_health_df.fetal_health)
```

```
X_fetal_matrix =
2126x21 Matrix{Float64}:
120.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  120.0  137.0  121.0  73.0  1.0
132.0  0.006  0.0  0.006  0.003  0.0  ...  1.0  141.0  136.0  140.0  12.0  0.0
133.0  0.003  0.0  0.008  0.003  0.0  ...  1.0  141.0  135.0  138.0  13.0  0.0
134.0  0.003  0.0  0.008  0.003  0.0  ...  0.0  137.0  134.0  137.0  13.0  1.0
132.0  0.007  0.0  0.008  0.0  0.0  ...  0.0  137.0  136.0  138.0  11.0  1.0
134.0  0.001  0.0  0.01  0.009  0.0  ...  3.0  76.0  107.0  107.0  170.0  0.0
134.0  0.001  0.0  0.013  0.008  0.0  ...  3.0  71.0  107.0  106.0  215.0  0.0
⋮
140.0  0.0  0.0  0.005  0.001  0.0  ...  0.0  145.0  143.0  145.0  2.0  0.0
140.0  0.0  0.0  0.007  0.0  0.0  ...  0.0  153.0  150.0  152.0  2.0  0.0
140.0  0.001  0.0  0.007  0.0  0.0  ...  0.0  152.0  148.0  151.0  3.0  1.0
140.0  0.001  0.0  0.007  0.0  0.0  ...  0.0  153.0  148.0  152.0  4.0  1.0
140.0  0.001  0.0  0.006  0.0  0.0  ...  0.0  152.0  147.0  151.0  4.0  1.0
142.0  0.002  0.002  0.008  0.0  0.0  ...  1.0  145.0  143.0  145.0  1.0  0.0

# Convert to numeric matrix
X_fetal_matrix = Matrix{Float64}(X_fetal)
```

Explanation:

- Separate predictors (X_fetal) and target (y_fetal) and convert target to categorical for classification.

10. Split data into training and testing for classification

```
X_fetal_table =
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	seve
1	120.0	0.0	0.0	0.0	0.0	0.0
2	132.0	0.006	0.0	0.006	0.003	0.0
3	133.0	0.003	0.0	0.008	0.003	0.0
4	134.0	0.003	0.0	0.008	0.003	0.0
5	132.0	0.007	0.0	0.008	0.0	0.0
6	134.0	0.001	0.0	0.01	0.009	0.0
7	134.0	0.001	0.0	0.013	0.008	0.0
8	122.0	0.0	0.0	0.0	0.0	0.0
9	122.0	0.0	0.0	0.002	0.0	0.0
10	122.0	0.0	0.0	0.003	0.0	0.0
⋮	more					

```

X_fetal_table = MLJBase.table(X_fetal_matrix, names=propertynames(X_fetal))

▶ ([279, 1880, 57, 905, 1114, 965, 165, 1350, 1028, ... more ,915], [398, 1709, 399, 706, 1805, 1
```

Train-test split

```
train, test = partition(eachindex(y_fetal), 0.7, shuffle=true, rng=123)
```

Explanation:

- Then convert predictors into an MLJ table
- Randomly split the data into 70% training and 30% testing for classification tasks.

11. Load and train a logistic classifier

```
LogisticClassifier
  • @load LogisticClassifier pkg=MLJLinearModels

  ⓘ For silent loading, specify `verbosity=0`.

  ⓘ import MLJLinearModels ✓

model_fetal = LogisticClassifier(
  lambda = 2.220446049250313e-16,
  gamma = 0.0,
  penalty = :l2,
  fit_intercept = true,
  penalize_intercept = false,
  scale_penalty_with_samples = true,
  solver = nothing)

  • model_fetal = LogisticClassifier()

mach_fetal =
  trained Machine; caches model-specific representations of data
  model: LogisticClassifier(lambda = 2.220446049250313e-16, ...)
  args:
    1: Source @604 ⇨ ScientificTypesBase.Table{AbstractVector{ScientificTypesBase.Continu
    2: Source @979 ⇨ AbstractVector{ScientificTypesBase.Multiclass{3}}

  • mach_fetal = machine(model_fetal, X_fetal_table, y_fetal) |> fit!

  ⓘ Training machine(LogisticClassifier(lambda = 2.220446049250313e-16, ...), ...).

  ⓘ Solver: MLJLinearModels.LBFGS{Optim.Options{Float64, Nothing}, @NamedTuple{}}
    optim_options: Optim.Options{Float64, Nothing}
    lbfgs_options: @NamedTuple{} NamedTuple()

y_pred =
  CategoricalArrays.CategoricalVector{Int64, UInt32, Int64, CategoricalArrays.CategoricalValue{I

  • # Evaluate
  • y_pred = predict_mode(mach_fetal, rows=test)

0.8589341692789969

  • begin
  •   import StatisticalMeasures ✓ : accuracy
  •   acc = accuracy(y_pred, y_fetal[test]) # This now returns a plain Float64
  • end

  • # Correct printing method
  • println("Classification Accuracy: ", round(acc * 100, digits=2), "%")

  ⓘ Classification Accuracy: 85.89%
```

@load LogisticClassifier pkg=MLJLinearModels

model_fetal = LogisticClassifier()

mach_fetal = machine(model_fetal, X_fetal_table, y_fetal) |> fit!

y_pred = predict_mode(mach_fetal, rows=test)

```

begin

  import StatisticalMeasures: accuracy

  acc = accuracy(y_pred, y_fetal[test])

end

# Print the classification accuracy

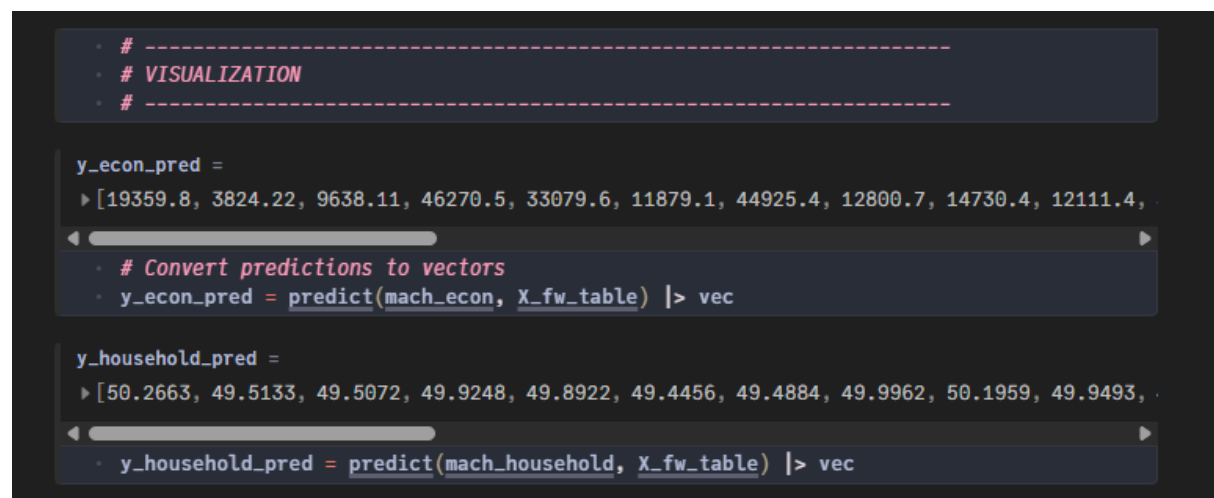
println("Classification Accuracy: ", round(acc * 100, digits=2), "%")

```

Explanation:

- Load a logistic regression model and fit it on the training data to classify fetal health status.
- Predicts classes using the model (`predict_mode`) and computes accuracy on the test set.
- Outputs classification performance.

12. Predict values for regression models



```

# -----
# VISUALIZATION
# -----

y_econ_pred =
  ▶ [19359.8, 3824.22, 9638.11, 46270.5, 33079.6, 11879.1, 44925.4, 12800.7, 14730.4, 12111.4, .
  ◀

# Convert predictions to vectors
y_econ_pred = predict(mach_econ, X_fw_table) |> vec

y_household_pred =
  ▶ [50.2663, 49.5133, 49.5072, 49.9248, 49.8922, 49.4456, 49.4884, 49.9962, 50.1959, 49.9493, .
  ◀

y_household_pred = predict(mach_household, X_fw_table) |> vec

```

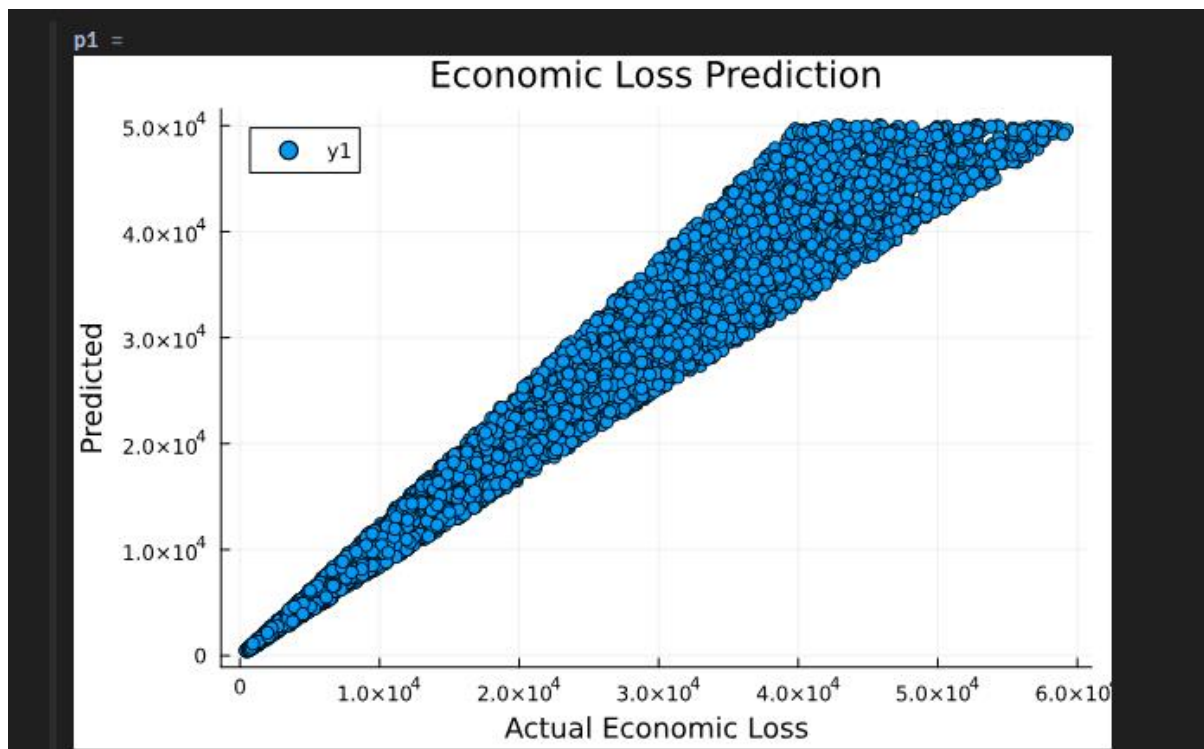
```
y_econ_pred = predict(mach_econ, X_fw_table) |> vec
```

```
y_household_pred = predict(mach_household, X_fw_table) |> vec
```

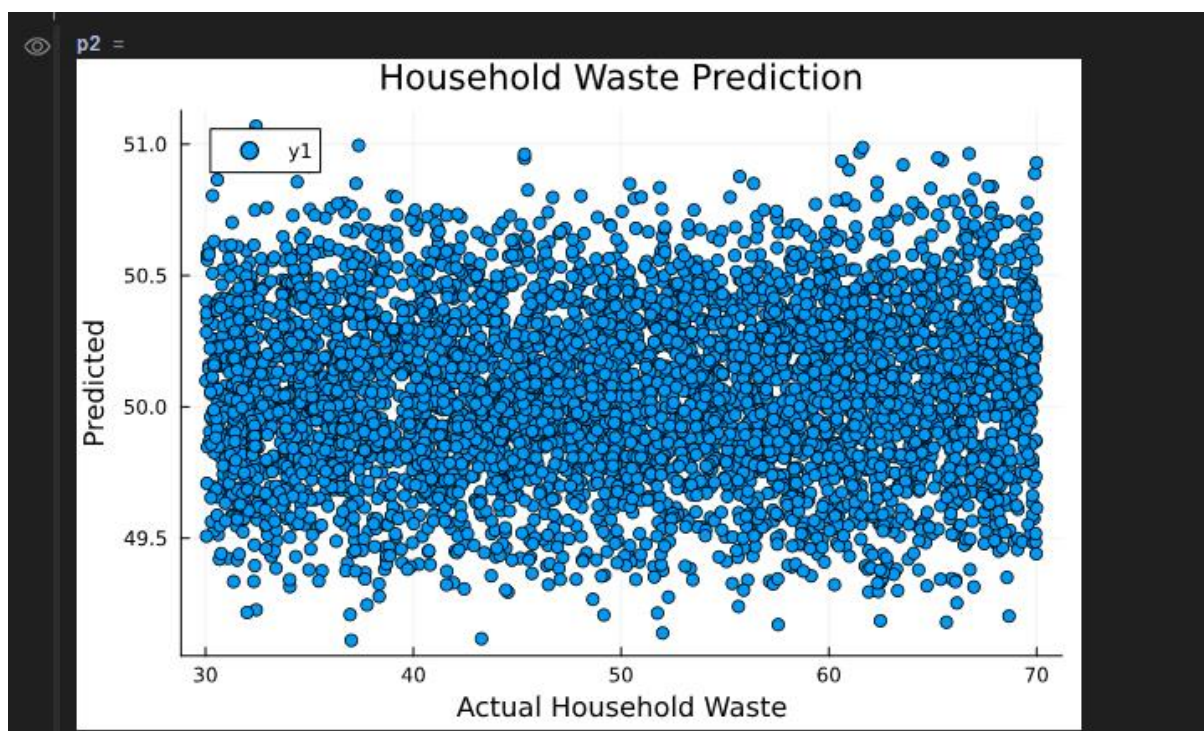
Explanation:

- Predict continuous values using the trained regression models and flatten the results to vectors.

13. Visualize regression predictions

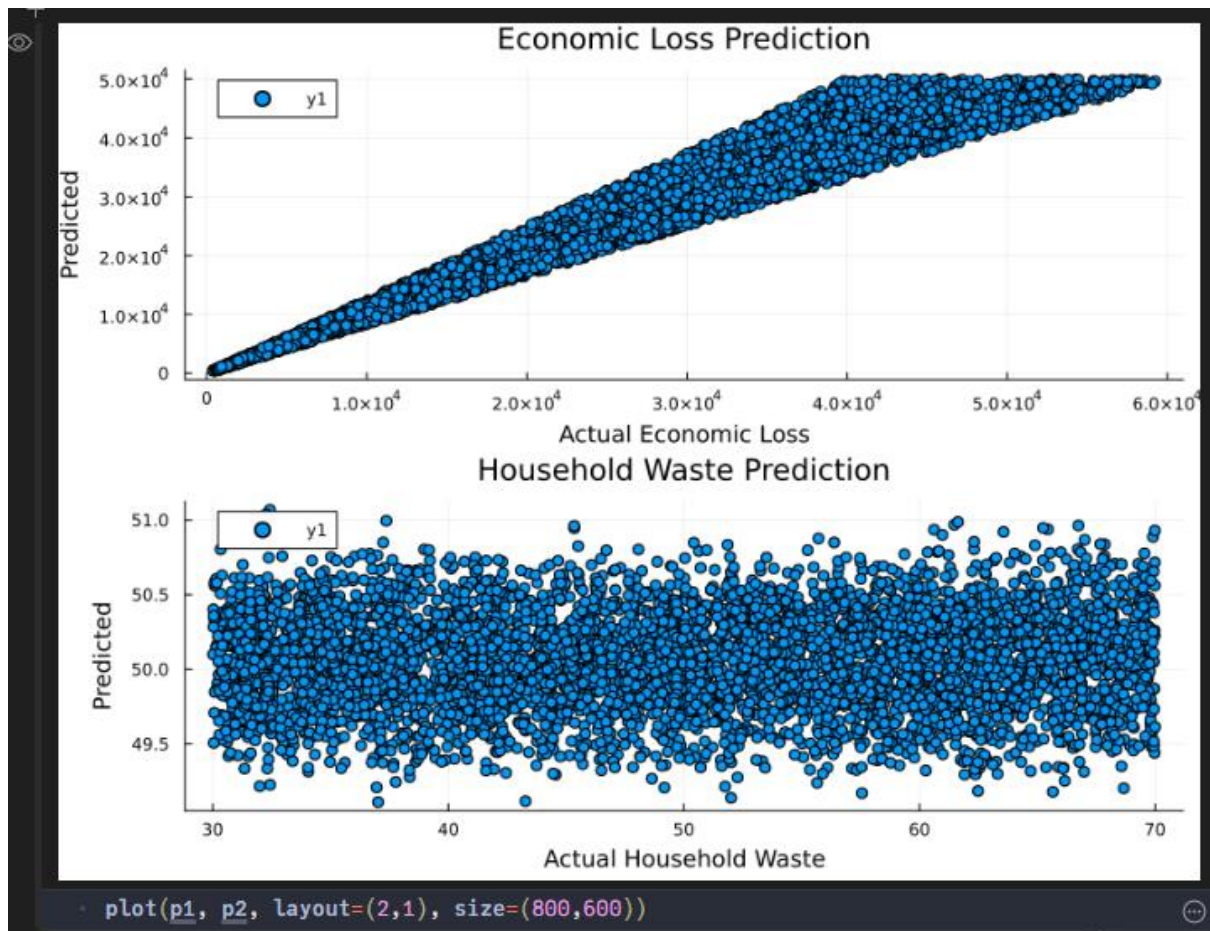


```
p1 = scatter(y_econ, y_econ_pred,  
            xlabel="Actual Economic Loss", ylabel="Predicted",  
            title="Economic Loss Prediction")
```



```
p2 = scatter(y_household, y_household_pred,
```

```
xlabel="Actual Household Waste", ylabel="Predicted",  
title="Household Waste Prediction")
```

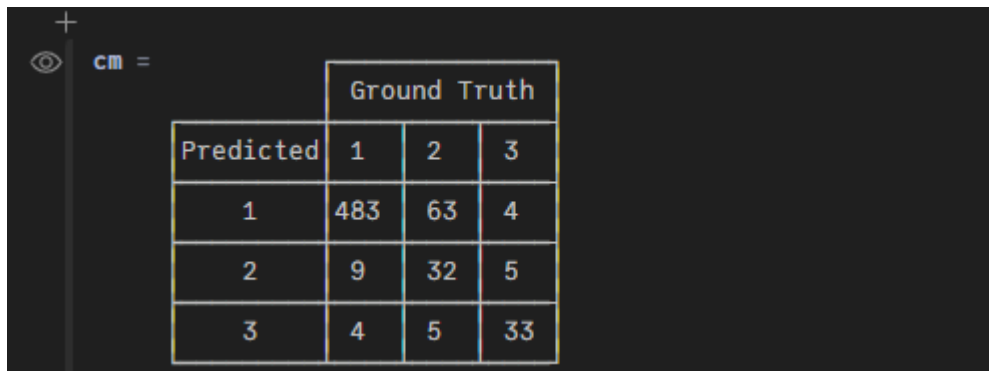


```
plot(p1, p2, layout=(2,1), size=(800,600))
```

Explanation:

- Creates scatter plots comparing actual vs predicted values for both regression tasks.

14. Generate confusion matrix for classification evaluation



Predicted	Ground Truth		
	1	2	3
1	483	63	4
2	9	32	5
3	4	5	33

```
cm = confusion_matrix(y_pred, y_fetal[test])
```

Explanation:

- Displays the confusion matrix, showing how well the classifier distinguishes between classes.

Short Summary

The Above code performs three machine learning tasks:

1. **Regression** - On the food wastage dataset to predict economic loss and household waste.
2. **Classification** - On fetal health data using logistic regression.
3. **Evaluation and visualization** - Using accuracy and scatter plots.